

JULY 2015		
3	4	5
10	11	12
17	18	19
24	25	26
31		
F	S	S

AUGUST 2015		
1	2	31/36
8	9	32
15	16	33
22	23	34
29	30	35
T	F	S
W	T	S

2015

JUL
THURSDAY
211-154 WK 31

30

Assignment - 4 (OS)

- (b) Distributed Deadlock
- S1: $P_1 \rightarrow P_2, P_3 \rightarrow P_4$
- S2: $P_2 \rightarrow P_5, P_5 \rightarrow P_6$
- S3: $P_6 \rightarrow P_1$

a) Combine all edges from the three sites:
 $P_1 \rightarrow P_2, P_2 \rightarrow P_5, P_5 \rightarrow P_6, P_6 \rightarrow P_1, P_3 \rightarrow P_4$
 Processes P_3 & P_4 form an unrelated waiting pair, they are not part of the other chain.

b) Deadlock detection:

Yes, deadlock exists here - there is a cycle:

$$P_1 \rightarrow P_2 \rightarrow P_5 \rightarrow P_6 \rightarrow P_1$$

This is a directed cycle, so processes involved in the deadlock are: P_1, P_2, P_5, P_6

P_3 and P_4 aren't part of the deadlock (they form a separate single wait edge $P_3 \rightarrow P_4$ only)

(7) Distributed 0.3

a) Probability local $\approx 1 - 0.3 = 0.7$

$$0.7 \times 5\text{ms} \approx 3.5\text{ms}$$

Probability remote ≈ 0.3

$$0.3 \times 25\text{ms} \approx 7.5\text{ms}$$

Sum: $3.5\text{ms} + 7.5\text{ms} \approx 11.0\text{ms}$

Expected access time $\approx 11\text{ms}$.

b) Strategy: Client-side read cache w/ server

31

JUL
FRIDAY
Wk 31 212-153

2015

27						
28	6	7	1	2	3	JULY
29	13	14	8	9	10	
30	20	21	15	16	17	
31	27	28	22	23	24	
			30	31	25	
Wk	M	T	W	T	F	S

7
14
21
28
M

- Leases (LRU replacement).
- Why → caching turns remote accesses into fast local reads cutting avg latency.
- Leases let clients safely use cached data w/o contacting the server on every read.
- LRU is a simple effective eviction policy.

Q8) a) Proposed Schedule (10 second window)

Do one full checkpoint at the start (1, 0):

200ms

Do incremental checkpoints every 1 sec afterwards at $t \geq 1, 2, 3, \dots, 9$ ($9 \text{ incremental} \times 50\text{ms} = 450\text{ms}$)

Total checkpoint overhead in 10 sec = $200 + 450 = 650 \text{ ms}$

$450 \approx 650 \text{ ms}$

Avg overhead = $650\text{ms} / 10\text{s} = 65\text{ms}$ per sec.

b) Reason:

$\rightarrow RPO \geq 1 \text{ Sec} \geq$ at most 1 sec of lost work.

Doing an incremental every 1 sec guarantees that in the worst case, we can recover within 1 sec.

→ Incremental checkpoints are cheap, so using them frequently is the low cost way to meet RPO.

→ Full Checkpoints are expensive so minimizing their frequency. One day full at the start provides the required baseline for replay.

1	2	3	4	5
6	7	8	9	10
11	12	13	14	15
16	17	18	19	20
21	22	23	24	25
26	27	28	29	30
T	W	T	F	S

1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30	31	1	2	3	4
T	W	T	F	S	S	S

2015

AUG
SATURDAY

213-152 WK 31

01

the chain of incrementals.

→ This min therefore satisfies the RPO while minimizing total checkpoints time in the 10sec window.

a) a) Sudden traffic spikes create hotspots and overload specific servers.

→ Maintaining low latency while balancing load across distant regions.

→ Keeping session/cart state local while still distributing load.

→ Avoiding cascading failures when one region becomes saturated.

Suitable algorithm = Power of Two-choice (P2C) load balancing bcoz it spreads load evenly and reduces hot spots.

b) → For low RPO:

Replicate critical data to other region continuously. Asynch replication gives small RPO; synchronous replication gives near-zero RPO.

→ For low RTO:

Use a global load balancer that quickly reroutes traffic when a region fails, plus automated fail over to promote replication.