

Assignment: Capstone

PART (A)

- 1) Modern systems rely on OS because they provide abstractions that hide hardware complexity & expose clean interfaces for programs

Key OS abstractions

- Process abstraction
virtual CPU, scheduling, isolation
- Memory abstraction
virtual memory, address space
- I/O abstraction
device independence, buffering, drivers

Scenarios

Process management

PCB creation, context switching, CPU Scheduling

Memory management

paging, segmentation, allocation, protection

I/O management

unified driver model, interrupt handling, DMA

why still needed today?

Even with advanced hardware, OS manages concurrency, security, resource allocation & fault tolerance which hardware alone can't do.

	Feature	Pros	Cons
2) Structure	Entire OS in one large kernel	Fast, efficient	Hard to maintain, low reliability
Monolithic	OS split into layers each using lower layer	Modular, easier debugging	Performance overhead
Layered	Minimal kernel, services run in user mode	High reliability, fault isolation	Slower due to more IPC
Microkernel			

Best for distributed web applications

Microkernel - because it offers fault isolation, modularity, maintainability & high reliability essential for distributed & cloud based deployments

3) Threads vs Processes

Threads are lighter because

- share code/data, processes have separate spaces
- thread control block (TCB) is smaller than PCB
- context switch between threads avoids MMU updates → faster
- threads consume fewer resources than full process creation

But threads are not always better

require synchronization → risk of race conditions

shared memory → debugging harder

process isolation offers better security

Thread management is more efficient, but less safe than process mgmt

4) Memory allocation

First fit & Best fit

Processes : 12 MB, 18 MB, 6 MB

Blocks : 20 MB, 10 MB, 15 MB

First fit

12 MB → 20 MB (remaining 8 MB)

18 MB → (next block 10 MB → skip, 15 MB → skip).
can't be allocated

6 MB → 10 MB (4 MB remaining)

Fragmentation 8 MB + 4 MB + unused 15 MB

Best fit

12 MB → 15 MB (remaining 3 MB)

18 MB → 20 MB (remaining 2 MB)

6 MB → 10 MB (remaining 4 MB)

Fragmentation 3 MB + 2 MB + 4 MB (total smaller)

First fit → more efficient.

PART (B)

5) FCFS, SJF, RR scheduling

Assume processes

$P_1 \text{ BT} = 10$

$P_2 \text{ BT} = 4$

$P_3 \text{ BT} = 6$

Gantt charts

FCFS

P_1	P_2	P_3	P_1	P_2	P_3
SJF	80	10	14	20	

RR ($q=4$)

Ready

P ₁	P ₂	P ₃	P ₁	P ₃	P ₁
----------------	----------------	----------------	----------------	----------------	----------------

Running

P ₁	P ₂	P ₃	P ₁	P ₃	P ₁
0	4	8	12	16	20

Average Time

Algo	Avg Waiting	Avg Turnaround
FCFS	High	High
SJF	lowest	lowest
RR	balanced	moderate

Best overall

Round Robin → balances throughput + fairness in multiprogram

6) Banker's Algo

(i) checks if resources requests keeps the system in a safe state
Allocation only if :-

Available = Need

Prevents circular wait → avoids deadlock.

b) Detection & Recovery approach

Use wait for graph

Detect cycles → indicates deadlock

Recovery methods

abort one/few transactions

roll back to checkpoint

release held locks

7) Producer-consumer using Semaphores

python

import threading, time from threading import

buffer = []

mutex = Semaphore(1)

empty = Semaphore(5)

full = Semaphore(0)

def producer():

 while True:

 item = "data"

 empty.acquire()

 mutex.acquire()

 buffer.append(item)

 mutex.release()

 full.release()

def consumer():

 while True:

 full.acquire()

 mutex.acquire()

 item = buffer.pop(0)

 mutex.release()

 empty.release()

Ensures

mutual exclusion → mutex

bounded buffer control → empty, full

no race conditions

8) Page replacement FIFO & LRU

Sequence : 2, 1, 4, 2, 3, 4, 3

FIFO

y_3	4	4	4	4	4
y_2	1	1	1	1	1
y_1	2	2	2	3	3
*	*	X	H	H	H

Page fault = 4

LRU

y_3	4	4	4	4	4
y_2	1	1	1	3	3
y_1	2	2	2	2	2
*	*	*	H	H	H

Page fault = 4

Frames differ after step 5 (FIFO replaced 2, LRU replaced 1) but total faults tie for this case

9) Distributed file systems
Two critical issues (concise)

- 10. Consistency vs Performance (Caching & Replication)
 - Cached / replicated copies improve read performance & fault tolerance but risk stale reads & write conflicts.
 - Choosing strong vs eventual consistency affects latency & correctness.

- 20. Metadata scalability & availability
 - Centralized metadata (file namespace, mapping to chunks) is a single point of failure or bottleneck;
 - distributing metadata increases consistency (consistency, fault-tolerance) but is necessary for global scale

b) Architectural approaches (concise)

1. Separation of metadata & data (GFS-style)

- small, replicated metadata service (masters) & many chunk / data servers. Clients contact metadata service for locations, then interact directly with chunk servers for data - reduces master load & improves throughput

2. Lease-based client caching + primary-backup

- give clients short leases for cached copies; primary leader handles writes & invalidates / updates to ensure consistency with bounded staleness.

3. Consistent hashing + sharding for scale

- partition namespace across metadata shards to distribute load; use replication within shards for availability.

4. Yield consistency

- choose file-level policies (strong for transactional files, eventual for logs / analytics) to balance latency & reliability.

10)

Synchronous checkpointing & recovery

Mechanism (pseudocode + ASCII diagram)

Coordinator: initiate checkpoint (k)

broadcast ("checkpoint-request", k)

wait for ACK's from all processes
if all ACKs received:

commit global checkpoint k to stable store
broadcast ("checkpoint-committed", k)

Process i on "checkpoint request k ":

flush outgoing messages to peers (or ensure deterministic logging)

write local state i_k to stable storage

send ACK to coordinator

ASCII timeline

Coordinator $\rightarrow P_1$: checkpoint-request

Coordinator $\rightarrow P_2$: checkpoint-request

P_1 : flush & write state $1_k \rightarrow$ ACK

P_2 : flush & write state $2_k \rightarrow$ ACK

Coordinator : receive all ACKs \rightarrow checkpoint committed

Recovery (concise)

On failure, restore each process from last committed global checkpoint k & replay logged messages if message logging is used

Strengths & Weakness

Synchronous

→ strength : simple, yields a consistent global state, easy recovery

→ weakness : High runtime overhead, stops progress (blocking) during checkpoint; poor scalability

Asynchronous

→ strength : low runtime interference; scalable

→ weaknesses: harder recovery (may require complex message logging / analysis); risk of domino effect unless message logging or coordinated techniques used.

- ii) Smart Home IoT (centralized OS)
- Scheduling Strategy (Algorithm + justification)
 - Hybrid & Immediate (Interrupt Handling + Preempt)
 - Priority Scheduling (EDF for soft-real time)
 - Priority-based with strict interrupt handling

Concretely:

- Hardware / firmware interrupts for security sensor (camera motion) preempt CPU immediately & trigger short, high-priority ISR that records event & signals a high-priority real-time task.
 - Scheduler: Use preemptive priority scheduler where security-processing threads have highest priority bounded execution time (or use EDF when deadlines explicit). Lower-priority tasks (lightning) run when no higher-priority tasks pending.
- Justification: Security events require minimal latency bounded response, preemptive priority + ISR ensures immediate attention & bounded interferences. EDF can be used if tasks have hard deadlines & utilization is analyzable.

Design notes:

- Keeps ISR's minimal (signal / enqueue) to avoid long interrupt handling in kernel.
- Offload heavy processing to dedicated edge AI (GPU) if available to reduce scheduling pressure on main OS.

b) IPC methods suitable (brief + reason)

- Message queues (local): decoupled, priority-aware useful for event notification (motion detected enqueue alert) easy to throttle and drop

low priority messages if overloaded.

2. Shared memory + mutex / semaphore : for high - bandwidth local data (camera frames) where copying is expensive. Use ring buffers + sequence numbers to avoid copying & reduce latency.
3. MQTT / lightweight broker (network) : for high - bandwidth local data (camera frames) where copying for device - to - cloud or cross room messaging - lightweight support QoS levels & is standard in IoT.
4. Unix domain sockets / RPC (local) : for structured requests between system components with lower overhead than TCP. Why these : Messages queues + shared memory combine low latency, controlled synchronization & decoupling - suitable for mixed - criticality IoT environments.

12) Case study : LINUX - Python demo of basic system calls I ran is short Python demo that uses POSIX - like system call wrappers (os.open, os.write, os.read) & also created threads (which use OS threading primitive). Output from running the script is below & file was saved at /mnt/data/os_syscall_demo.txt

code run

```
import os, threading, time
fd = os.open ("/mnt/data/os_syscall_demo.txt", os.O_CREAT | os.O_WRONLY, 0644)
os.write (fd, b"OS syscall demo : file created using os."
open / os.write \n line 2\n")
```

os.close(fd)

fd = os.open("/mnt/data/os-syscall-demo.txt",
R00NLY)

content = os.read(fd, 1024)

os.close(fd)

print(content.decode())

def worker(n):

print(f"Worker {n} is running (thread id: {the-
get_ident()})")

time.sleep(0.1)

print(f"Worker {n} is done")

t1 = threading.Thread(target=worker, args=(1,))

t2 = threading.Thread(target=worker, args=(2,))

t1.start(); t2.start()

t1.join(); t2.join()

Observed output

OS syscall demo: file created using os.open/os.write
Line 2\n

worker 1 running (threading id: 1392651156374)

worker 2 running (threading id: 139265107211)

worker 1 done

worker 2 done

2 threads completed

Note: Python threading calls into OS threading API
C p threads on POSIX